SRISHTI M

AI Product Manager

CS Engineering Graduate | Next Leap Top PM Fellow | Bengaluru, India

UX Design • AI Development • Web Development • NLP • Computer Vision

PORTFOLIO

GenAI • LLM Products • NLP • Computer Vision • Product Strategy

2025

Table of Contents

About Me & Core Skills 3
Positioning Statement 4
Case Study 01 — AI Meeting Briefing Generator 5
Case Study 02 — WhatsApp Memory Retrieval System 9
Case Study 03 — ChatGPT Voice India Private Mode PRD 13
Case Study 04 — AI Face Recognition System (CV/ML) 17
Why Hire Srishti as an AI PM 20

About Me

I am a Computer Science Engineering graduate and a Top PM Fellow at Next Leap, passionate about building AI-powered products that solve real user problems. My unique advantage is a full-stack background spanning UX design, web development, NLP, and Computer Vision — enabling me to bridge user empathy with technical depth.

I specialize in GenAI and LLM product strategy, combining research-first thinking with strong execution instincts. I am actively seeking AI PM roles in Bengaluru where I can shape the next generation of intelligent products.

Quick Facts

🎓 CS Engineering Graduate
🏆 Top PM Fellow — Next Leap
📍 Bengaluru, India
🎯 Target: AI PM (GenAI/LLM)
💼 Experience: Fresher (0–1 yr)

Core Skills & Tools

Domain	Skills & Tools
AI Development	NLP, Computer Vision, PCA, ANN, ML pipelines, Python, scikit-learn, OpenCV
Product Management	PRD Writing, User Research, RICE/ICE, Funnel Analysis, Roadmapping, A/B Testing
Design	UX Design, Wireframing, Figma, Prototyping, User Flows
Development	Web Design & Development, React, HTML/CSS, JavaScript
Analytics	Metrics Definition, North Star Metrics, Funnel Analysis, KPI Design
AI Automation	AI integrations, workflow automation, LLM API integrations (GPT, Whisper ASR)

Positioning Statement

Why Srishti for an AI PM Role?

A rare combination of CS engineering depth, hands-on AI development (Computer Vision + NLP), UX design instincts, and product management training — making her one of the few fresher candidates who can truly bridge user needs, engineering realities, and business outcomes in AI products.

The Unfair Advantage

Skill Area	PM Superpower	Evidence
UX Design	Deep user empathy + rapid prototyping	Built full UI/UX for ChatGPT PRD — Figma + Lovable prototype
AI Development	Speaks engineer's language; understands model tradeoffs	Built PCA + ANN face recognition system; NLP pipelines
Web Dev	Reduces dev-PM communication gap	End-to-end web implementation experience
NLP/CV Projects	Hands-on AI product intuition	450-sample face recognition dataset, 74.34% accuracy at k=100
Next Leap Fellowship	Structured PM thinking & case methodology	4 polished product case studies across AI + consumer products

CASE STUDY 01 — Next Leap Fellowship Project

AI Meeting Briefing Generator

Proactively preparing knowledge workers for meetings using AI context synthesis

Role	Type	Users	Tools
Product Manager (Solo)	GenAI / Productivity	Knowledge Workers — PMs, Sales, Founders	LLM APIs, Calendar APIs, Figma

The Problem

Problem Statement

Knowledge workers attend 8–12 meetings per week. Preparation is manual and fragmented — jumping across email, calendar, notes, and attendee profiles. Meetings begin with context reconstruction, wasting time and weakening outcomes.

PM Insight: Why This Problem, Why Now

- LLMs now enable pre-meeting context synthesis that was technically impossible 2 years ago
- Calendar and email APIs are mature, enabling automated trigger-based workflows
- The largest user pain is not during the meeting — it's the 15–30 minutes of scattered prep before it

User Research & Insights

Primary Research

- Survey conducted with 32 respondents; 6 in-depth interviews
- Primary segment: Busy knowledge workers — PMs, Sales, Founders with high meeting frequency

Key Personas

Persona	Context	JTBD (Job to Be Done)
Product Manager	Frequent internal + external syncs; scattered context across tools	When a meeting is starting, help me recall scattered context so I sound prepared
Sales Professional	Client-facing; needs deal history, stakeholder profiles quickly	Before a client call, help me focus on what matters most without digging through email

Solution: AI Meeting Briefing

Core Solution

An AI-powered briefing system that triggers 4 hours before each meeting, synthesizes context

from email, calendar, past meetings, and public profiles, and delivers a structured briefing via email/dashboard — with a feedback loop for continuous improvement.

System Architecture

- **Calendar detects upcoming meetings (T-4 hours)** Trigger Layer:
- **Email service, Calendar service, Meeting history, LinkedIn profiles** Data Collection:
- **Context aggregation → Key theme extraction → Info prioritization → Summarization** AI Intelligence Layer:
- **Email delivery + Dashboard delivery** Output Layer:
- **User signals feed back into quality improvement model** Feedback Loop:

User Flow

- Upcoming meeting detected → AI briefing auto-generated → User reviews briefing → User provides feedback → Meeting preparedness confirmed

MVP Scope & Prioritization

Feature	Priority	RICE Rationale
Upcoming meetings view	P0	Core entry point — no value without this
AI meeting briefing generation	P0	Primary value delivery
Suggested focus areas	P1	Increases meeting effectiveness
Feedback capture	P1	Enables learning loop — critical for AI quality
Confirmation state	P2	Trust building and habit formation

Why AI Briefings Over Alternatives

Solution	Verdict	Reason
Manual Notes	Rejected	High effort, limited contextual relevance, user-dependent
Static Summaries	Rejected	Low context, misses real-time meeting-specific signals
AI Briefings	Selected	Proactive, multi-source, learns via feedback — RICE: High/High/Med-High/Med

Metrics & Success Criteria

Metric Type	Metric	Target
North Star	% of meetings where briefing is opened & used	> 50% within 90 days
Engagement	Briefing open rate	> 60%
Quality	Time spent reviewing briefing	> 90 seconds avg

Retention	Repeat usage rate (weekly)	> 40%
Outcome	Self-reported prep confidence score	> 4/5 rating

Funnel Analysis

Funnel Stage	Conversion Rate	PM Insight
Eligible meetings detected	100%	Baseline — calendar integration working
AI briefings generated	85%	15% failure rate in generation — model/API issue
Briefing opened	47%	Critical drop: discovery & notification are key challenges
Briefing reviewed	25%	Habit formation needed — progressive onboarding opportunity
Meeting preparedness improved	8%	North Star conversion — low but expected for v1 MVP

Risks, Guardrails & Future Scope

Guardrails

- Privacy & consent: User-controlled data access, explicit opt-in for each data source
- Hallucination checks: Briefing quality scoring + explainability for each insight
- User control: Manual override and opt-out at any time

Risk Mitigation

- **Active feedback loop retrains ranking model** Low briefing relevance →
- **Full transparency on data sources used per briefing** Trust erosion →
- **Human review mode for sensitive meeting types** Over-automation →

Future Scope

- Personalization: Briefing adapts based on user role and meeting history patterns
- Voice briefings: Audio delivery for commute-time consumption
- CRM + task integrations: Pull Salesforce, Jira, Notion context into briefings

CASE STUDY 02 — Next Leap Fellowship Project

WhatsApp Memory Retrieval System

Turning WhatsApp from a chat app into a trusted memory layer through contextual recall

Role	Type	Platform	Market
Product Manager (Solo)	NLP / Consumer AI	WhatsApp Mobile (3B+ MAU)	~$2.4B Annual Revenue; $70B projected by 2034

The Problem

Problem Statement

WhatsApp is increasingly used as an informal memory system, but current search assumes keyword recall. Human memory works on contextual cues — who sent it, roughly when, what type of content. This mismatch causes failed searches, manual scrolling, and the socially costly act of re-asking others to resend information.

Root Cause Analysis

- Search depends on exact keyword recall — not how human memory actually works
- No cross-chat contextual indexing exists in WhatsApp today
- Cost of retrieval exceeds cost of re-asking → users give up
- Low confidence in search results → trust erosion → feature abandonment

Research & Insights

Quantitative Signals

- Majority of users fail to find information on first search attempt
- High retry and abandonment rates in existing search flows
- Manual scrolling is the dominant fallback behavior
- If Recall reaches 5% of daily searches and reduces failures by 30% → ~2.5% retention lift

Qualitative Insights

- Users remember WHO sent something, not the exact words
- Search is a last resort, not first instinct
- Re-asking feels socially inefficient — users experience embarrassment
- Trust in search drops sharply after 1–2 failed attempts

Personas

Persona	Profile	Pain Point	JTBD
Pranay, 30 — PM	Uses WhatsApp for work coordination; receives links,	Loses trust after 1–2 failed search attempts;	When urgent, retrieve important details using

	docs, addresses across many chats	anxious in urgent moments	partial context — who sent it, when, what type
Sejal, 25 — Software Engineer	Part of many groups; WhatsApp is her default memory store; never organizes upfront	Overwhelmed by volume; resorts to endless scrolling or re-asking	When info floods in, I want WhatsApp to surface what I need later without manual organization

Solution: Contextual Recall

Core Solution

A memory-first retrieval experience where users express partial memory context (who sent it, roughly when, what type of content) to surface the most relevant past messages across all chats — ranked by contextual importance, not keyword match.

Why Contextual Recall vs. Alternatives

Solution	RICE Score	Verdict
Contextual Recall	Highest (High/High/High/Med)	Selected — matches how users remember; works with metadata; privacy-safe
Manual Organization (Tags/Bookmarks)	Moderate	Rejected — requires upfront discipline; fails retroactively
Auto Chat Summaries	Low (Low confidence)	Rejected — privacy risk under E2EE; hard to verify accuracy

6-Screen User Flow

Screen	Feature	Design Rationale
1	Recall Entry Point	Dedicated tab — signals users can retrieve without exact keywords
2	Contextual Filters	Express partial memory: Who? When? What type?
3	Ranked Recall Results	Cross-chat results ranked by contextual importance
4	No Confidence Match	Honest transparency — builds trust by not showing wrong results
5	Result Explanation	"Why this result" — reduces trial-and-error behavior
6	Jump to Original	Navigates to exact message in original chat context

Prototype Validation Results

- 62% found first recall result relevant on first try
- 78% completed the recall flow end-to-end
- 41% used contextual filters to narrow results

System Design: Privacy-Preserving Architecture

Key Principle: Metadata-only intelligence — no breaking of E2EE

- Content Extraction: Raw message metadata only (sender, timestamp, type)
- Importance Scoring: Relevance ranking without reading encrypted content
- Recall Index: Fast cross-chat retrieval index built on metadata signals
- Privacy Guardrails: Enforce privacy at every layer; category suppression for sensitive chats
- Feedback Loop: User signals (open, dismiss, retry) improve ranking over time silently

Metrics Framework

Level	Metric	Definition	Why It Matters
North Star	Successful Recall Completion Rate	% of recall sessions where user opens the correct original message	Direct measure of user value delivered
L0 — Adoption	Recall Feature Usage Rate	% of eligible users using contextual recall vs. manual scroll	Measures discoverability and habit formation
L0 — Quality	First Result Open Rate	% of recall attempts where the first result is opened	Measures ranking quality and user trust
L0 — Efficiency	Time to Successful Recall	Average seconds from recall start to correct result opened	Measures urgency-moment effectiveness
L1 — Retention	7-Day Repeat Recall Usage	% of users who use recall again within 7 days	Measures long-term trust and value
Guardrail	Privacy Complaints Rate	Any spike triggers immediate investigation	Mandatory for E2EE trust maintenance

Risk & Mitigation Matrix

Risk	Impact	Mitigation Strategy
Recall results feel creepy/invasive	Trust erosion	Metadata-only processing + explicit 'Why this result' explanation on every result
Wrong result ranked #1	User loses confidence	Conservative ranking + instant re-rank on dismiss signal
Sensitive chats surfaced	Privacy breach	Category suppression + on-device classification for sensitive content
Feature slower than manual scroll	Feature abandonment	On-device recall index + cached recent contexts for <500ms response
Users don't discover recall	Low adoption	First-time education tooltip + contextual nudge at search failure moments

CASE STUDY 03 — PRD / Next Leap Fellowship

ChatGPT Voice India — Private Mode

A privacy-first voice interaction mode for ChatGPT that enables natural voice use in public spaces

Roles	Type	Target	Status
PM + Designer + Developer (Srishti M)	GenAI / Voice UX PRD	Mobile users 18–35 in India	In Review — Launch: Nov 2025

Problem Definition

Core Problem

68% of ChatGPT users avoid voice input in shared spaces (public transport, cafes, offices) due to privacy discomfort. 42% cite 'I don't want others to hear the response' as the top blocker. Voice usage is highest at night (home) and lowest during commute hours — indicating discomfort, not usability failure.

Why Now

- Voice interfaces are growing rapidly; ChatGPT needs a privacy-safe voice mode to remain competitive
- Alexa, Google Assistant, and WhatsApp have all introduced privacy-aware audio modes
- India's commute culture creates massive unmet need — millions of daily transit users who could benefit

Solution: Auto-Detect Private Mode for Voice

How It Works

Trigger	Action	User Experience
Earphones connected	Private Mode auto-activates	Toast notification: 'Private Mode On — Responses will be text-only'
Voice input detected	Whisper ASR transcribes speech	Waveform animation + transcript shown ('Got it!')
Response generated	Text-only output, no audio	Text response displayed; 'Play Audio' optionally available
Earphones removed	Mode auto-deactivates	Toast: 'Earphones removed — switching back to audio mode'

System Architecture (4 Layers)

- **Mic input detection, earphone connection events, OS audio routing** Device Layer:
- **Whisper ASR converts speech → text (< 2.5s end-to-end latency target)** AI Layer:

- **Decides output mode (audio vs. text-only) based on earphone state + user setting** Logic Layer:
- **Displays text response + privacy badge, mutes speaker output, shows trust cues** UI Layer:

Goals & Success Metrics

Priority	Goal	Success Metric	Target
P0	Increase voice feature adoption	% uplift in active voice users	+25% lift
P1	Improve comfort with public voice use	Self-reported safety score	+30% increase
P2	Maintain fast performance	End-to-end response latency	< 2.5 seconds
P3	Reduce mode confusion	Mode-switch abandonment rate	< 5%

UX Design Decisions

UX Element	Design Choice	Rationale
Discovery	Banner inside Voice tab: 'Speak privately, get text-only replies'	Contextual placement drives activation at the right moment
Trust signal	Lock icon + earphone icon + purple accent color	Visual language reinforces privacy without requiring reading
Mode confirmation	'Private Mode On' toast + persistent status bar indicator	Reduces anxiety — user always knows current mode
Graceful failure	'No Confident Match' state with refinement options	Maintains trust even when ASR fails; avoids frustration

Edge Cases & Error Handling

Scenario	Expected Behavior
Earphones disconnect mid-reply	Toast: 'Earphones removed — switching back to audio mode'
Both Bluetooth + wired connected	Prioritize most recently active audio device
User disables Bluetooth permissions	Manual Private Mode toggle available in settings
ASR fails or poor connection	'We couldn't hear you — try again or type instead'
Bluetooth speaker connected	Treated as public — Private Mode does not activate

Launch Plan

Phase	Date	Owner	Action

Prototype	28 Oct 2025	Srishti M	Figma prototype + Lovable interactive demo ready
Dogfooding	30 Oct – 3 Nov	Internal	Internal testing and feedback collection
Beta Rollout (10%)	6 Nov 2025	Mobile PM	A/B test vs. control users without feature
Full Rollout	15 Nov 2025	Engineering + Ops	Complete rollout with App Store release note

Non-Goals (Explicit Scope Decisions)

- No speech-to-speech output in v1 — descoped to v2
- No desktop/web support — mobile-only for MVP
- No multilingual voice support — English-only for MVP
- No whisper/low-volume voice output — future iteration
- Not solving dictation/transcription use cases — different user need

CASE STUDY 04 — AI Development Project

AI Face Recognition System

Built a production-grade face recognition pipeline using PCA dimensionality reduction and Artificial Neural Networks

Role	Type	Stack	Outcome
AI Developer / ML Engineer	Computer Vision — Classification	Python, scikit-learn, PCA, ANN, OpenCV, NumPy	74.34% accuracy at optimal k=100 eigenfaces

Project Overview & Problem

Core Challenge

High-dimensional facial image data (90,000 features per image from 300×300 pixel grayscale images) makes direct classification computationally infeasible. The goal was to build an efficient, accurate face recognition system that scales — identifying individuals from a 450-sample, 450-class dataset.

Technical Architecture

Pipeline: 5-Stage ML Workflow

Stage	Process	Technical Detail
1 — Data Preparation	Image preprocessing	Grayscale conversion, resize to 300×300, normalize pixel values, flatten to 90,000-dim vector
2 — Dimensionality Reduction	PCA (Eigenfaces)	Extracted top-k principal components; reduced 90,000 features to k=100 eigenfaces for optimal performance
3 — Feature Projection	Eigenface projection	Projected all train/test images onto eigenface orthonormal basis — train: (337, 150), test: (113, 150)
4 — Classification	MLP Classifier (ANN)	Multi-layer perceptron with Adam optimizer; hidden layers (50,); max_iter=1000; cross-validation tuning
5 — Evaluation	Accuracy analysis	k-sweep experiment across k=10, 20, 50, 100, 150 components to find optimal trade-off

Key Experimental Results

PCA Components (k)	Classification Accuracy	Interpretation
k = 10	47.79%	Insufficient features — underfitting; key facial

		variance lost
k = 20	58.41%	Improving — basic facial structure captured
k = 50	62.83%	Good generalization — mid-level features retained
k = 100	74.34% (PEAK)	Optimal sweet spot — maximum accuracy before overfitting begins
k = 150	68.14%	Slight drop — noise introduced as k exceeds optimal range

PM Translation of Results

As k increases, accuracy rises — until diminishing returns kick in at k=100. This is the classic PM tradeoff: more data/complexity doesn't always mean better outcomes. The 'optimal k' insight directly maps to product decisions like feature bloat vs. focused scope.

PM-Lens: Learnings Applied to Product Thinking

Technical Learning	PM Application
Model accuracy vs. computational cost tradeoff	Speed vs. accuracy tradeoff in AI product features — when 'good enough' is better than perfect
Cross-validation prevents overfitting	A/B testing prevents premature product decisions based on insufficient data
PCA component selection (k-sweep)	Feature prioritization — not all features add value; ruthless scoping improves outcomes
Feedback loop via prediction probability scores	AI products need confidence signals + explainability to build user trust
Dataset size limits generalization	AI product quality is bounded by training data — data strategy is a PM responsibility

Technical Skills Demonstrated

- Machine Learning: PCA, LDA, MLP Classifier, cross-validation, hyperparameter tuning
- Computer Vision: Image preprocessing pipeline, eigenfaces, feature extraction
- Python ecosystem: scikit-learn, NumPy, OpenCV, matplotlib
- ML evaluation: Accuracy metrics, confusion analysis, train/test split strategy
- AI ethics awareness: Identified bias risks in face recognition, recommended mitigation strategies

Why Hire Srishti as an AI PM?

What Companies Want in an AI PM	What Srishti Brings
Deep technical understanding of ML/AI	Built NLP pipelines, PCA+ANN systems, AI integrations — not just theoretical knowledge
User empathy and UX thinking	UX designer background; created full user flows, prototypes, personas across every project
Product rigor: PRDs, metrics, roadmaps	4 complete case studies with RICE prioritization, North Star metrics, funnel analysis
GenAI / LLM product sense	Designed AI features leveraging LLMs, Whisper ASR, and context synthesis architectures
Communication across functions	CS + UX + PM — speaks fluently to engineers, designers, and business stakeholders
Research-first thinking	Primary research (surveys, interviews) backs every product decision across all projects

Target Roles & Companies in Bengaluru

Company Type	Companies	Why Srishti Fits
Indian AI Startups	Sarvam AI, Krutrim, Haptik, Leena AI	NLP + GenAI background directly relevant; startup agility mindset
Conversational AI	Haptik, Uniphore, Yellow.ai	NLP projects + UX design = ideal combo for voice/chat AI products
B2B AI SaaS	Darwinbox, Zoho, Freshworks	PRD + metric skills; understands enterprise user needs
Big Tech AI Teams	Google, Microsoft, Atlassian	Technical depth + structured PM thinking required — strong match
Consumer Fintech AI	Razorpay, CRED, PhonePe, Meesho	Consumer empathy + AI integration skills

STAR Stories Summary

Competency	Story Summary	Project
Data-Driven Decisions	Ran k-sweep experiment (k=10 to k=150) to identify optimal PCA components; chose k=100 based on accuracy peak — not intuition	Face Recognition
User Research Leadership	Conducted 32-person survey + 6 in-depth interviews; surfaced insight that 'recall failure happens in high-pressure moments' — shaped entire WhatsApp product strategy	WhatsApp Recall
Cross-Functional Execution	Acted as PM + Designer + Developer on ChatGPT PRD; created full Figma prototype,	ChatGPT Voice PRD

	system architecture, and launch plan solo	
Handling Ambiguity	Designed privacy-preserving AI architecture (metadata-only, E2EE-safe) for WhatsApp Recall — navigating real technical + regulatory constraints	WhatsApp Recall
Stakeholder Communication	Translated NLP/CV technical findings into business impact narratives: '5% recall adoption → 2.5% retention lift → millions of hours saved'	All Projects

Ready to Build AI Products That Matter.

Let's connect — srishtikushil.6@gmail.com

Bengaluru, India | Open to AI PM roles in Product, GenAI, NLP, and LLM domains